

```

import datetime
from packaging import version
from collections import Counter
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# TensorFlow and tf.keras
import tensorflow as tf
from tensorflow import keras
import tensorflow_datasets as tfds

%matplotlib inline
np.set_printoptions(precision=3, suppress=True)

def plot_graphs(history, metric):
    plt.plot(history.history[metric])
    plt.plot(history.history['val_'+metric], '')
    plt.xlabel("Epochs")
    plt.ylabel(metric)
    plt.legend([metric, 'val_'+metric])

#register ag_news_subset so that tfds.load doesn't generate a checksum (mismatch)
!python -m tensorflow_datasets.scripts.download_and_prepare --register_checksums --

# https://www.tensorflow.org/datasets/splits
# The full `train` and `test` splits, interleaved together.
ri = tfds.core.ReadInstruction('train') + tfds.core.ReadInstruction('test')
dataset_all, info = tfds.load('ag_news_subset', with_info=True, split=ri, as_supervised)

2021-08-09 17:08:57.834956: I tensorflow/stream_executor/platform/default/dso_loader.cc:44] Successfully opened dynamic library libcudart.so.10.1
I0809 17:09:00.013056 139638028969856 download_and_prepare.py:200] Running download_and_prepare.py
ag_news_subset
2021-08-09 17:09:00.020596: I tensorflow/core/platform/cloud/google_auth_provider.cc:150] default_credentials: Creating default credentials from environment.
2021-08-09 17:09:00.128816: I tensorflow/core/platform/cloud/google_auth_provider.cc:150] default_credentials: Creating default credentials from environment.
2021-08-09 17:09:00.208541: I tensorflow/core/platform/cloud/google_auth_provider.cc:150] default_credentials: Creating default credentials from environment.
I0809 17:09:00.282195 139638028969856 dataset_info.py:434] Load pre-computed dataset info for ag_news_subset
2021-08-09 17:09:00.296452: I tensorflow/core/platform/cloud/google_auth_provider.cc:150] default_credentials: Creating default credentials from environment.
2021-08-09 17:09:00.504638: I tensorflow/core/platform/cloud/google_auth_provider.cc:150] default_credentials: Creating default credentials from environment.
I0809 17:09:00.721039 139638028969856 dataset_info.py:361] Load dataset info for ag_news_subset
I0809 17:09:00.722968 139638028969856 download_and_prepare.py:138] download_and_prepare.py
I0809 17:09:00.723343 139638028969856 dataset_builder.py:357] Generating dataset builder for ag_news_subset
Downloading and preparing dataset ag_news_subset/1.0.0 (download: 11.24 MiB, generated: 1.00 KiB)
2021-08-09 17:09:00.865821: I tensorflow/core/platform/cloud/google_auth_provider.cc:150] default_credentials: Creating default credentials from environment.
2021-08-09 17:09:00.969605: I tensorflow/core/platform/cloud/google_auth_provider.cc:150] default_credentials: Creating default credentials from environment.

```

✓ 0s completed at 1:03 PM



Dl Size....: 0 MiB [00:08, ? MiB/s]

Extraction completed....: 0 file [00:08, ? file/s]

Dl Completed....: 0% 0/1 [00:08<?, ? url/s]

Dl Size....: 1 MiB [00:08, 8.72s/ MiB]

Dl Completed....: 0% 0/1 [00:08<?, ? url/s]

Dl Size....: 2 MiB [00:08, 8.72s/ MiB]

Dl Completed....: 0% 0/1 [00:08<?, ? url/s]

Dl Size....: 3 MiB [00:08, 8.72s/ MiB]

Dl Completed....: 0% 0/1 [00:08<?, ? url/s]

Dl Size....: 4 MiB [00:08, 8.72s/ MiB]

Dl Completed....: 0% 0/1 [00:08<?, ? url/s]

Dl Size....: 5 MiB [00:08, 8.72s/ MiB]

Dl Completed....: 0% 0/1 [00:08<?, ? url/s]

Dl Size....: 6 MiB [00:08, 8.72s/ MiB]

Dl Completed....: 0% 0/1 [00:08<?, ? url/s]

Dl Size....: 7 MiB [00:08, 8.72s/ MiB]

Dl Completed....: 0% 0/1 [00:08<?, ? url/s]

Dl Size....: 8 MiB [00:08, 8.72s/ MiB]

Dl Completed....: 0% 0/1 [00:08<?, ? url/s]

Dl Size....: 9 MiB [00:08, 8.72s/ MiB]

Dl Completed....: 0% 0/1 [00:08<?, ? url/s]

Dl Size....: 10 MiB [00:08, 8.72s/ MiB]

Dl Completed....: 0% 0/1 [00:08<?, ? url/s]

Dl Size....: 11 MiB [00:08, 8.72s/ MiB]

EDA

```
# classes dictionary
categories =dict(enumerate(info.features["label"].names))
categories
train_categories = [categories[label] for label in dataset_all.map(lambda text, lat
Counter(train_categories).most_common())
```

```
doc_sizes = []
corpus = []
for example, _ in dataset_all.as_numpy_iterator():
    enc_example = encoder(example)
    doc_sizes.append(len(enc_example))
    corpus+=list(enc_example.numpy())
```

```
# register ag_news_subset so that tfds.load doesn't generate a checksum (mismatch)
!python -m tensorflow_datasets.scripts.download_and_prepare --register checksums --
```

```
# dataset, info = tfds.load('ag_news_subset', with_info=True, split=['train[:]', 'train[:95%]', 'train[:114000]'], as_supervised=True)
```

```
# train dataset, test dataset = dataset['train'],dataset['test']
```

1

```
test_dataset = test_dataset.batch(BATCH_SIZE).prefetch(tf.data.experimental.AUTOTUNE)
```

```
for example, label in train_dataset.take(2):
    print('texts: ', example.numpy()[:3])
    print()
    print('labels: ', label.numpy()[:3])
```

```
texts: [b'The wreckage of a Detroit case about a "sleeper operational combat
b'NOVEMBER 29, 2004 (IDG NEWS SERVICE) - Malicious hackers have compromised '
b'Russian Svetlana Kuznetsova came back twice from a break down in the second
```

```
labels: [0 3 1]
texts: [b'Gossip-seekers eager to dig up dirt on David Beckham risk ending up
b'NEW YORK -- New York boxing officials have suspended former heavyweight cha
b'Reuters - The trial of Indonesian\\cleric Abu Bakar Bashir over charges re'
```

```
labels: [3 1 0]
```

```
VOCAB_SIZE=1000
```

```
encoder = tf.keras.layers.experimental.preprocessing.TextVectorization(
    max_tokens=VOCAB_SIZE)
encoder.adapt(train_dataset.map(lambda text, label: text))
```

```
num_classes = 4
```

```
model = tf.keras.Sequential([
    encoder
    ,tf.keras.layers.Embedding(input_dim=len(encoder.get_vocabulary()),
                              output_dim=64
                              # Use masking to handle the variable sequence lengths
                              ,mask_zero=True)
    ,tf.keras.layers.Bidirectional([tf.keras.layers.LSTM(64, return_sequences=True),
    ,tf.keras.layers.LSTM(64, return_sequences=False)]
    ,tf.keras.layers.GlobalAveragePooling1D()
    ,tf.keras.layers.Dense(num_classes))])
```

```
Epoch 6/200  
1782/1782 [=====] - 30s 17ms/step - loss: 0.3827 - a  
Epoch 7/200  
1782/1782 [=====] - 32s 18ms/step - loss: 0.3777 - a  
Epoch 8/200  
1782/1782 [=====] - 32s 18ms/step - loss: 0.3736 - a
```

```
test_loss, test_acc = model.evaluate(test_dataset)
```



```
print('Test Loss: {}'.format(test_loss))  
print('Test Accuracy: {}'.format(test_acc))
```

```
119/119 [=====] - 1s 10ms/step - loss: 0.4034 - accu  
Test Loss: 0.4034428000450134  
Test Accuracy: 0.8501315712928772
```

```
history_dict = history.history  
history_dict.keys()
```

```
history_df=pd.DataFrame(history_dict)  
history_df.tail(10).round(3)
```

```
plt.show()
```



```
preds = model.predict(test_dataset)
```

```
preds.shape
```

```
, loss=tf.keras.losses.sparse_categorical_crossentropy() # if we set this  
, metrics=['accuracy'])
```

Simple RNN significantly increases the amount of time needed to train the model for approximately the same performance

```
history = model.fit(train_dataset  
                    , epochs = 200  
                    , validation_data=validation_dataset  
                    , callbacks=[tf.keras.callbacks.EarlyStopping(monitor='val_accu
```

1	0.400	0.850	0.410	0.844
2	0.421	0.851	0.437	0.847
3	0.399	0.858	0.404	0.859
4	0.387	0.862	0.405	0.859
5	0.376	0.864	0.393	0.860

